

In []:

```
%pip install spark
%pip install pyspark

from pyspark.sql import SparkSession

spark = SparkSession.builder \
    .appName("example") \
    .getOrCreate()
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>
 Requirement already satisfied: spark in /usr/local/lib/python3.9/dist-packages (0.2.1)
 Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>
 Requirement already satisfied: pyspark in /usr/local/lib/python3.9/dist-packages (3.3.2)
 Requirement already satisfied: py4j==0.10.9.5 in /usr/local/lib/python3.9/dist-packages (from pyspark) (0.10.9.5)

In []:

```
# DataFrames
# Use the DataFrames in Apache Spark to calculate the following:
# 1. Fatalities per month (12.5 pts): Calculate the fatalities per month. Show
# with the most fatalities, in descending order.
# 2. Fatalities per year (12.5 pts): Calculate the fatalities per year. Show
# the most fatalities, in descending order.

from pyspark.sql.functions import year, month, countDistinct

fatalities_df = spark.read.format("csv").option("header", True).load("combine

# Create new columns 'fatality_month' from 'FAT YEARMONTH' column
fatalities_df = fatalities_df.withColumn('fatality_month', fatalities_df['FAT

# Perform aggregation by month and count the distinct fatality IDs
fatalities_by_month = fatalities_df.groupby('fatality_month').agg(countDistinct

# Create new columns 'fatality_year'
fatalities_df = fatalities_df.withColumn('fatality_year', substring('FAT_YEAR

# Perform aggregation by year and count the distinct fatality IDs
fatalities_by_year = fatalities_df.groupby('fatality_year') \
    .agg(countDistinct('FATALITY_ID').alias('f
    .orderBy('fatalities', ascending=False)

# Show the results
fatalities_by_year.show(10)
fatalities_by_month.show()
```

+-----+	
fatality_year	fatalities
+-----+	
2005	1453
2011	1335
2018	1050
2021	984
1999	906
2020	901
2008	825
2017	775
2019	732
2007	712
+-----+	

only showing top 10 rows

+-----+	
fatality_month	fatalities
+-----+	
07	3381
08	3248
06	1955
04	1697
05	1634
+-----+	

In []:

```
# SparkSQL
# Use SparkSQL to calculate the following:
# 1. Fatalities per month (12.5 pts): Calculate the fatalities per month. Show
# with the most fatalities, in descending order.
# 2. Fatalities per year (12.5 pts): Calculate the fatalities per year. Show
# the most fatalities, in descending order.

from pyspark.sql import SparkSession

# Create SparkSession
spark = SparkSession.builder.appName("fatalities_analysis").getOrCreate()

# Load fatalities data from CSV file
fatalities_df = spark.read.format("csv").option("header", True).load("combine")

# Register DataFrame as temporary view
fatalities_df.createOrReplaceTempView("fatalities")

# Perform SQL query to calculate fatalities by year
fatalities_by_year = spark.sql("SELECT SUBSTR(FAT_YEARMONTH, 1, 4) as fatality_year \
                                FROM fatalities \
                                GROUP BY fatality_year \
                                ORDER BY fatalities DESC")

# Create a temporary view for SparkSQL queries
fatalities_df.createOrReplaceTempView("fatalities")

# Perform the SQL query to calculate fatalities by month
fatalities_by_month = spark.sql("SELECT SUBSTRING(FAT_YEARMONTH, 5, 2) AS fatality_month \
                                  FROM fatalities \
                                  GROUP BY fatality_month \
                                  ORDER BY fatalities DESC")

# Show the result
fatalities_by_year.show(10)
fatalities_by_month.show()
```

+-----+	
fatality_year	fatalities
+-----+	
2005	1453
2011	1335
2018	1050
2021	984
1999	906
2020	901
2008	825
2017	775
2019	732
2007	712
+-----+	

only showing top 10 rows

+-----+	
fatality_month	fatalities
+-----+	
07	3381
08	3248
06	1955
04	1697
05	1634
+-----+	